

CORE-ET CPU Subsystem Micro- Architecture Reference

Table of Contents

Document Control	1
Revision History	1
Chapter Status	1
Preface	2
Purpose	2
Audience	2
Sources and Source of Truth	2
Conventions	2
1. ET-SoC-1 Overview	3
1.1. ET-SoC-1 at a Glance	3
1.2. The ET-Minion Hierarchy	3
1.2.1. ET-Minion Core	4
1.2.2. Neighborhood	5
1.2.3. Minion Shire	6
1.2.4. Minion Shire Array	7
1.3. ET-Maxion and Service Processor	7
1.3.1. ET-Maxion Cores	7
1.3.2. Service Processor	8
1.4. Mesh Network-on-Chip	8
1.4.1. Mesh Topology and Mesh Stops	9
1.4.2. Memory Shires	9
1.4.3. PCIe Shire	10
1.4.4. I/O Shire	10
2. ET-Minion Core	11
2.1. Core Overview	11
2.1.1. Block Partitioning	11
2.1.2. Hardware Threads	11
2.1.3. Instruction Flow Through the Core	11
2.1.4. Interfaces to the Neighborhood	11
2.1.5. Clocks, Resets and Power	11
2.1.6. Debug Support	11
2.2. Front End	12
2.2.1. Overview	12
2.2.2. Interface	13
2.2.3. Pipeline Stages	17
2.2.4. Interaction with the L0 Micro-Cache	18
2.2.5. Thread Buffer	22
2.2.6. Entry Invalidation, Redirect and Thread Start	26

2.2.7. Decoders	28
2.2.8. Instruction FIFO	29
2.2.9. Thread Scheduler	30
2.2.10. Halt and Debug Program Buffer	31
2.3. Integer Pipeline	32
2.3.1. Overview	32
2.3.2. Pipeline Stages	32
2.3.3. Interfaces	32
2.3.4. Register File and Operand Read	32
2.3.5. Immediate Generation and ALU	32
2.3.6. Multiply/Divide Unit	32
2.3.7. Scoreboards and Hazards	32
2.3.8. Hart Interleaving and Stalls	32
2.3.9. Branches, Redirects and Replay	32
2.3.10. Interface to the Data Cache	32
2.3.11. Interface to the VPU	32
2.3.12. CSRs and Privilege Modes	33
2.3.13. Exceptions and Interrupts	33
2.3.14. Debug and Breakpoints	33
2.3.15. Performance Counters	33
2.3.16. Verification Notes	33
2.4. Data Cache	33
2.4.1. Overview	33
2.4.2. Organization and Operating Modes	33
2.4.3. Interfaces	33
2.4.4. Request Pipeline	33
2.4.5. Address Translation and PMA Checks	33
2.4.6. Tag, Data and LRU Arrays	33
2.4.7. Miss Handling	33
2.4.8. Replay Queue	33
2.4.9. Writeback and Store Merge	33
2.4.10. Atomic Operations	33
2.4.11. Cache-Management Operations	33
2.4.12. Scratchpad Control	33
2.4.13. Tensor Load and Reduce	34
2.4.14. Interface to the Neighborhood	34
2.4.15. Errors	34
2.4.16. Performance	34
2.4.17. Verification Notes	34
2.5. Vector Processing Unit	34
2.5.1. Overview	34

2.5.2. Interfaces	34
2.5.3. Control and Micro-Instruction Decode	34
2.5.4. Register Files	34
2.5.5. Lanes	34
2.5.6. Mask Unit	34
2.5.7. Bypass Network	34
2.5.8. Shift and Swizzle Unit	34
2.5.9. TxFMA Unit	34
2.5.10. Integer Multiply-Accumulate Unit	34
2.5.11. Transcendental Unit	34
2.5.12. Tensor Operations	34
2.5.13. Interfaces to the Integer Pipeline and Data Cache	35
2.5.14. Performance	35
2.5.15. Verification Notes	35
3. Neighborhood	36
3.1. Neighborhood Overview	36
3.1.1. Partitioning	36
3.1.2. Shared Resources and Core Connections	36
3.1.3. Voltage-Domain Crossings	36
3.1.4. External Interfaces	36
3.2. Front End – I-Cache Interface	36
3.2.1. Overview	36
3.2.2. Interface Signals	36
3.2.3. Request Arbitration	36
3.2.4. Request and Response Timing	36
3.2.5. Miss, Sleep and Fill-Done Protocol	36
3.2.6. Fault and Error Reporting	36
3.2.7. Livelock Avoidance	36
3.2.8. Known Limitations	36
3.3. L0 Micro-Cache	36
3.3.1. Overview	37
3.3.2. Internal Blocks	37
3.3.3. Pipeline Stages	37
3.3.4. Instruction TLB	37
3.3.5. Micro-Tag and Micro-Data Arrays	37
3.3.6. PMA Check	37
3.3.7. Replacement	37
3.3.8. Miss Handling	37
3.3.9. L1 Bypass	37
3.3.10. Software Prefetch	37
3.4. L1 Instruction Cache	37

3.4.1. Overview	37
3.4.2. Organization	37
3.4.3. Tag, LRU and Data Arrays	37
3.4.4. External Data RAM	37
3.4.5. Request Arbitration Between L0 Micro-Caches	37
3.4.6. Miss Path	38
3.4.7. Errors and ECC	38
3.5. Shared Page-Table Walker	38
3.5.1. Clients and Arbitration	38
3.5.2. Walk Sequence and Page Sizes	38
3.5.3. PTE Cache and Invalidation	38
3.6. Memory Request and Response Path	38
3.6.1. Miss, Evict and Fill FIFOs	38
3.6.2. Request Arbitration	38
3.6.3. Response Routing and Backpressure	38
3.7. Tensor Support	38
3.7.1. Cooperative Tensor Load	38
3.7.2. Tensor Store Buffer	38
3.7.3. TBOX Router	38
3.8. Synchronization, Messaging, Registers and PMU	38
3.8.1. Fast Local Barrier	38
3.8.2. Local Message Network	38
3.8.3. IPI and FCC Paths	38
3.8.4. Neighborhood Registers and APB Access	38
3.8.5. Debug Channel	38
3.8.6. Performance Monitoring Unit	39

Document Control

Revision History

Chapter Status

Preface

Purpose

Audience

Sources and Source of Truth

Conventions

1. ET-SoC-1 Overview

1.1. ET-SoC-1 at a Glance

ET-SoC-1 is a many-core RISC-V system-on-chip developed primarily for artificial-intelligence and machine-learning (AI/ML) workloads. It is designed as a flexible, fast and efficient inference engine for ML applications. Its compute resources are arranged in a multi-level hierarchy that delivers high sustained performance on regular, parallelizable workloads, such as those of ML algorithms.

The device integrates 1,093 general-purpose 64-bit RISC-V cores of three types:

- 1,088 ET-Minion cores for ML computation
- 4 ET-Maxion cores for management and general-purpose tasks
- 1 Service Processor for boot and device management

The cores are connected to on-die SRAM, external memory and a set of standard interfaces through a mesh network-on-chip (NoC). [Table 1](#) summarizes the device.

Table 1. ET-SoC-1 key characteristics

Item	Description
ET-Minion cores	1,088 in-order, dual-threaded RV64IMFC cores, each with an 8-lane vector/tensor unit
ET-Maxion cores	4 single-threaded, superscalar, out-of-order 64-bit RISC-V cores
Service Processor	1 single-threaded core derived from the ET-Minion
On-die SRAM	140 MB. Each 1 MB block is configurable as L2 cache, as part of the chip-wide L3 cache, or as scratchpad memory.
Interconnect	Mesh NoC with 44 mesh stops
External memory	LPDDR4X, 16 channels × 16 bits, up to 4,266 MT/s
PCI Express	Gen4, 8 lanes. Endpoint, root complex, or both at once.
Other interfaces	eMMC, USB 2.0 OTG, I2C/SMBus, I3C, SPI, UART, GPIO. Debug access through JTAG and a USB 2.0 device port.

1.2. The ET-Minion Hierarchy

The ET-Minion cores are organized in three levels, as shown in [Figure 1](#):

- **ET-Minion core**
- **Neighborhood**: 8 ET-Minion cores
- **Minion Shire**: 4 Neighborhoods (32 ET-Minion cores)

Each level adds resources that are shared by the units it contains. The device contains 34 Minion Shires, which together form the Minion Shire array ([Section 1.2.4](#)).




Figure 1. ET-Minion hierarchy

1.2.1. ET-Minion Core

The ET-Minion is a dual-threaded, in-order, single-issue 64-bit RISC-V core that implements RV64IMFC. Its main features are:

- Two hardware threads (harts)
- 4 KB private L1 data cache, partly configurable as scratchpad memory
- Vector processing unit (VPU) with eight identical lanes operating in lock-step. Each lane contains:
 - Floating-point multiply-add unit (FMA): one 32-bit or two 16-bit operations
 - Two integer multiply-add units (IMA): four 8-bit multiply-accumulates each
 - Integer unit (INT)
 - Transcendental unit (TRANS)
- Esperanto instruction extensions for vector, tensor, atomic, message, cache-management and fast-synchronization operations

The core consists of four top-level blocks, shown in [Figure 2](#):

- **Front end (FE):** fetches instructions from the shared instruction cache and decodes them
- **Integer pipeline**
- **L1 data cache and scratchpad**
- **VPU**




Figure 2. ET-Minion core

1.2.2. Neighborhood

A Neighborhood groups eight ET-Minion cores that share the following resources. In Figure 3, the L0 micro-caches are labeled "I-Cache".

- **L1 instruction cache:** 32 KB, organized as 128 sets $\times$ 4 ways with 64-byte lines. It is shared by all eight cores, and its data RAMs are located outside the Neighborhood.
- **L0 micro-caches:** two, each shared by four cores. Each holds 16 fully associative cache lines, sits between its cores and the L1 instruction cache, and accesses the L1 instruction cache on a miss.
- **Page-table walkers (PTWs):** two, each shared by four cores. Each PTW serves the data caches of its four cores and the instruction-cache side.
- **Request/response path:** a single path to the memory hierarchy outside the Neighborhood. Requests from the instruction and data caches are arbitrated onto it.
- **Performance monitoring unit (PMU):** twelve 64-bit event counters shared by the eight cores. Eight counters count ET-Minion events and four count Neighborhood events.




Figure 3. Neighborhood



The PTWs were not implemented in the first silicon revision (A0). The RTL described in this document implements them.

1.2.3. Minion Shire

A Minion Shire groups four Neighborhoods (Figure 4). It also contains:

- **Shire Cache:** 4 MB, in four 1 MB banks, connected to the Neighborhoods through a crossbar. Programmable ESRs divide the cache into three kinds of partition:
 - L2 cache, private to the Shire
 - Slice of the chip-wide distributed L3 cache
 - Scratchpad memory
- **Local services:** uncacheable accesses, fast inter-processor interrupts (IPIs), fast local barriers and fast local credit counters.
- **ESR banks:** the Esperanto System Registers (ESRs) configure the blocks in the Shire.
- **NoC interface:** L2 misses and uncacheable accesses leave the Shire here, bound for other Shires or the memory controllers.




Figure 4. Minion Shire

1.2.4. Minion Shire Array

The 34 Minion Shires, 1,088 ET-Minion cores in total, form the Minion Shire array. The roles of the Shires are assigned by software:

- **Compute:** 32 Minion Shires operate in parallel as the compute array.
- **Master (management):** one Minion Shire receives work and distributes it to the compute array.
- **Spare:** one Minion Shire.

All Minion Shires are identical in hardware, so any Minion Shire can take any of these roles.

1.3. ET-Maxion and Service Processor

In addition to the ET-Minion cores, the device contains two other core types. Both are located in the I/O Shire ([Section 1.4.4](#)). This document does not describe them beyond this section.

1.3.1. ET-Maxion Cores

The ET-Maxion is a single-threaded, superscalar, out-of-order 64-bit RISC-V core intended for management and general-purpose tasks. The device contains four identical ET-Maxion cores. Together with a 4 MB shared L2/L3 cache they form the Maxion Neighborhood, and a system bus and coherency hub in the I/O Shire connects them. [Figure 5](#) shows the ET-Maxion pipeline.




Figure 5. ET-Maxion pipeline

1.3.2. Service Processor

The Service Processor (SP) is a single-threaded ET-Minion core without the vector/tensor extensions. It has its own ROM and a 1 MB scratchpad SRAM, but none of the Neighborhood or Shire shared resources.

- **Boot.** The SP is the first major block released from reset. It executes the boot code from ROM and initializes the device interfaces, including DDR, PCIe and USB.
- **Run time.** It provides management services for the whole device.
- **Access.** The SP can access every resource on the device, and some resources are accessible only to it.

The I/O Shire also contains a hardware root of trust, and six secure mailboxes for communication between four agents:

- Service Processor
- Master Minion Shire
- ET-Maxions
- PCIe host

1.4. Mesh Network-on-Chip

The mesh NoC is the top-level interconnect of ET-SoC-1 and connects all Shires. Figure 6 shows the device with the Shire array on the mesh, the external memory and PCIe interfaces, and the peripheral, Service Processor and debug networks.




Figure 6. ET-SoC-1 block diagram

1.4.1. Mesh Topology and Mesh Stops

The mesh is organized as an 8×6 grid, and each Shire attaches to it at a mesh stop. The Memory Shires occupy the east and west edges, four on each side, which leaves the corners of the grid unoccupied. The result is 44 mesh stops, listed in [Table 2](#).

Table 2. Shires on the mesh

Shire	Count	Contents
Minion Shire	34	32 ET-Minion cores and a 4 MB Shire Cache
Memory Shire	8	LPDDR4X memory controller
PCIe Shire	1	PCIe controllers and PHY
I/O Shire	1	ET-Maxions, Service Processor and peripherals
Total	44	

1.4.2. Memory Shires

Each of the eight Memory Shires contains one 32-bit LPDDR4X controller that drives two 16-bit channels at up to 4,266 MT/s. The Memory Shires also support global atomic operations.

1.4.3. PCIe Shire

The PCIe Shire contains two dual-mode PCIe Gen4 controllers that share an 8-lane PHY. It supports three lane configurations:

- Single 8-lane endpoint or root complex
- Single 4-lane endpoint or root complex
- 4-lane endpoint together with a 4-lane root complex

The interface is normally factory-configured as an endpoint.

1.4.4. I/O Shire

The I/O Shire contains the ET-Maxions and the Service Processor ([Section 1.3](#)). It also contains the device peripherals, which are grouped on three networks according to which agents may access them.

- **Service Processor network.** Accessible only to the SP. It contains:
 - SPI, including the boot-flash controller
 - I2C, including the link to the external power-management controller
 - UART, including the SP console
 - GPIO, timers, watchdog and boot control
- **Peripheral network.** Shared by the SP, the ET-Maxions and the master Minion Shire. It contains:
 - eMMC, USB 2.0 OTG, I2C/SMBus, I3C, SPI and UART
 - GPIO, timers and watchdog
- **Debug network.** Provides debug access through JTAG and a USB 2.0 device port.

2. ET-Minion Core

2.1. Core Overview

2.1.1. Block Partitioning

Figure: ET-Minion core block diagram.

2.1.2. Hardware Threads

2.1.3. Instruction Flow Through the Core

The Front End fetches instructions for both threads from the Neighborhood L0 micro-cache, expands compressed instructions, decodes them, and presents one instruction per cycle to ID, the first stage of the integer pipeline. The integer pipeline executes the instructions in order and passes memory operations to the data cache and vector operations to the VPU. In the opposite direction, the integer pipeline redirects the Front End to a new fetch address after a branch misprediction, exception or other change of control flow, and stalls a thread that must not issue.




Figure 7. Front End in the ET-Minion core

2.1.4. Interfaces to the Neighborhood

Table: ET-Minion core interface groups.

2.1.5. Clocks, Resets and Power

2.1.6. Debug Support

2.2. Front End

2.2.1. Overview

The Front End is the interface between the Neighborhood L0 micro-cache and the integer pipeline. It holds the program counter of each thread, fetches instructions from the L0 micro-cache, and delivers them, expanded and decoded, to the ID stage of the integer pipeline at a rate of one instruction per cycle. It serves both threads of the core, and restarts fetch at a new address when the integer pipeline redirects a thread.

To hide the latency of the L0 micro-cache, which is shared by four ET-Minion cores, each thread uses a double buffer of two 256-bit entries, each holding half a cache line. A new block can therefore be requested while the thread consumes the instructions already received. Each entry has a dedicated register that holds the address of its block.

[Table 3](#) lists the submodules and [Figure 8](#) shows how they are connected.

Table 3. Front End submodules

Submodule	Instances	Function
Thread buffer (frontend_thread_buffer)	2, one per thread	64-byte instruction buffer and the thread's program counter. Issues fetch requests to the L0 micro-cache.
RVC expander (frontend_rvc_expander)	2, one per thread buffer	Expands compressed (RVC) instructions to 32 bits.
Request arbiter (arb_lru_data)	1	2:1 LRU arbiter. Selects the thread whose fetch request is sent to the L0 micro-cache.
Integer decoder (intpipe_decode)	1	Decodes the instruction into integer-pipeline control signals.
VPU decoder (vpu_decoder)	1	Decodes the instruction into VPU control signals.
Instruction FIFO	2, one per thread	Two-entry queue of decoded instructions waiting for the integer pipeline.
Thread scheduler (frontend_thread_sched)	1	Selects the thread whose instruction is presented to the integer pipeline.




Figure 8. Front End block diagram

2.2.2. Interface

2.2.2.1. I/O Signals

Table 4 lists the ports of `frontend_top`. Connected to gives the signal in `core_top`. A width of 2 denotes one bit per thread.

Table 4. Front End I/O signals

Signal	Dir	Width / type	Connected to	Description
System				
<code>clock</code>	in	1	<code>clock</code>	Core clock. Each thread buffer derives a gated clock from it.
<code>reset</code>	in	1	<code>reset_w</code>	Warm reset. Synchronous.
<code>reset_debug</code>	in	1	<code>reset_d</code>	Debug reset. Clears the program-buffer control and, while the thread is halted, refills the program buffer with <code>EBREAK</code> .
<code>chicken_bit</code>	in	1	<code>chicken_bit_frontend</code>	Disables Front End clock gating. Driven by the <code>min_frontend_clock_gate_disable</code> chicken bit.
Thread control				

Signal	Dir	Width / type	Connected to	Description
<code>f0_thread_enabled</code>	in	2	<code>enabled</code>	Thread enable. A rising edge starts fetch at the reset vector; a falling edge stops fetch.
<code>f0_reset_vector</code>	in	48	<code>reset_vector</code>	Fetch address after reset or thread enable.
<code>vm_status</code>	in	2 × <code>minion_vm_status</code>	<code>vm_status</code>	Virtual-memory status of each thread, forwarded with every fetch request.
Fetch request to the L0 micro-cache				
<code>f1_icache_req_valid</code>	out	1	<code>icache_req_valid</code>	The F1 request register holds a request.
<code>f1_icache_req</code>	out	<code>frontend_icache_req</code>	<code>icache_req</code>	Fetch request: thread, VM status and address.
<code>f1_icache_req_ready</code>	in	1	<code>icache_req_ready</code>	The Neighborhood accepts the request in this cycle.
Fetch response from the L0 micro-cache				
<code>f5_icache_resp_valid</code>	in	1	<code>icache_resp_valid</code>	A response is present. It carries no thread identifier.
<code>f5_icache_resp_miss</code>	in	1	<code>icache_resp_miss</code>	The response is a miss and carries no instruction data.
<code>f5_icache_resp</code>	in	<code>icache_frontend_resp</code>	<code>icache_resp</code>	32-byte block and status.
<code>f6_icache_fill_done</code>	in	1	<code>icache_fill_done</code>	The L0 micro-cache has completed a line or translation fill.
Redirect and stall from the integer pipeline				
<code>f0_core_req_valid</code>	in	2	<code>id_core_fe_req_valid</code>	Redirect. Flushes the thread and restarts fetch at <code>f0_core_req.pc</code> .
<code>f0_core_req</code>	in	2 × <code>minion_fe_req</code>	<code>id_core_fe_req</code>	Redirect PC, speculation status and debug-mode transition.
<code>id_core_stall</code>	in	2	<code>id_core_fe_stall</code>	Thread must not issue. Raised for WFI , a fence, exclusive mode held by the other thread, and integer or FP scoreboard stalls.
Instruction to the integer pipeline (ID)				
<code>id_inst_valid</code>	out	1	<code>id_fe_core_resp_valid</code>	An instruction is presented.

Signal	Dir	Width / type	Connected to	Description
<code>id_inst_ready</code>	in	1	<code>id_core_fe_resp_ready</code>	The integer pipeline accepts the instruction.
<code>id_inst_thread_id</code>	out	1	<code>id_fe_core_resp_thread_id</code>	Thread of the instruction.
<code>id_inst_data</code>	out	<code>frontend_core_resp</code>	<code>id_fe_core_resp</code>	Instruction, PC and fetch status.
<code>id_intpipe_ctrl</code>	out	<code>minion_control</code>	<code>id_intpipe_ctrl</code>	Integer decode.
<code>id_vpu_decoder_sigs</code>	out	<code>vpu_ctrl_sigs_t</code>	<code>id_vpu_decoder_sigs</code>	VPU decode. Valid only for VPU instructions.
<code>id_vpu_core_ctrl</code>	out	<code>vpu_minion_id_ctrl</code>	<code>id_frontend_vpu_ctrl</code>	VPU and mask register usage, for hazard checks in the integer pipeline.
Debug				
<code>halt</code>	in	2	<code>debug_in_core.halt</code>	A halt request is pending.
<code>halted</code>	in	2	<code>debug_out.halted</code>	The thread is in debug mode.
<code>debug_ffb_wdata</code>	in	64	<code>debug_ffb_wdata</code>	Program-buffer write data.
<code>debug_ffb_en</code>	in	4	<code>debug_ffb_en</code>	Program-buffer word write enables.
<code>debug_ffb_thread_sel</code>	in	1	<code>debug_ffb_thread_sel</code>	Thread whose program buffer is written.
<code>debug_ffb_exec</code>	in	2	<code>debug_ex_program_buffer</code>	Start executing the program buffer.

2.2.2.2. Data Structures

Table 5. `frontend_icache_req`: fetch request (58 bits)

Field	Width	Description
<code>thread_id</code>	1	Requesting thread.
<code>vm_status</code>	8	Virtual-memory status of the thread (Table 6).
<code>addr</code>	49	Fetch address. The L0 micro-cache returns the aligned 32-byte block that contains it.

Front End addresses are 49 bits wide: the 48-bit virtual address and one additional bit that allows a non-canonical 64-bit address to be represented.

Table 6. `minion_vm_status`: virtual-memory status (8 bits)

Field	Width	Description
<code>prv</code>	2	Current privilege level.
<code>mprv</code> , <code>mpp</code> , <code>sum</code> , <code>mrx</code>	1, 2, 1, 1	Corresponding <code>mstatus</code> fields.

Field	Width	Description
<code>debug</code>	1	Thread is in debug mode.

Table 7. `icache_frontend_resp`: fetch response (261 bits)

Field	Width	Description
<code>data</code>	256	32-byte instruction block.
<code>page_fault</code>	1	Instruction page fault.
<code>access_fault</code>	1	Instruction access fault.
<code>cacheable</code>	1	Block is in a cacheable region.
<code>bus_err</code>	1	Bus error on the line.
<code>ecc_err</code>	1	ECC error on the line.

Table 8. `minion_fe_req`: redirect (52 bits)

Field	Width	Description
<code>pc</code>	49	New fetch PC.
<code>speculative</code>	1	Thread has older instructions in the integer pipeline. Driven every cycle, independently of <code>f0_core_req_valid</code> . Used for replay (Section 2.2.5.6).
<code>debug_info.halt</code>	1	Redirect enters debug mode.
<code>debug_info.resume</code>	1	Redirect leaves debug mode.

Table 9. `frontend_core_resp`: instruction to the integer pipeline (89 bits)

Field	Width	Description
<code>pc</code>	49	PC of the instruction.
<code>inst_bits</code>	32	Instruction. Compressed instructions are delivered expanded.
<code>pf0, af0</code>	1 each	Page fault / access fault on the entry holding the start of the instruction.
<code>pf1, af1</code>	1 each	Page fault / access fault on the entry holding the upper half of an instruction that crosses entries.
<code>replay</code>	1	Instruction must be replayed.
<code>rvc</code>	1	Instruction was compressed.
<code>bus_error, ecc_error</code>	1 each	Bus error / ECC error on the entry or entries holding the instruction.

Table 10. `frontend_thread_data`: instruction FIFO entry (179 bits)

Field	Width	Description
<code>core_resp</code>	89	Instruction and fetch status (Table 9).

Field	Width	Description
<code>vpu_ctrl_sigs</code>	45	VPU decode (<code>vpu_ctrl_sigs_t</code>).
<code>intpipe_ctrl</code>	45	Integer decode (<code>minion_control</code>).

Table 11. Decoded control words

Structure	Width	Content
<code>minion_control</code>	45	Legal instruction; software-emulated (M-code) instruction; branch, <code>JAL</code> , <code>JALR</code> ; integer source reads and destination write; ALU operand selects, function and width; immediate format; memory operation, command and size; FP/VPU destination write; mask register read/write; multiply/divide; CSR command; fence; global memory access; gather/scatter; graphics; implicit use of <code>x31</code> .
<code>vpu_ctrl_sigs_t</code>	45	VPU function; execution unit (TxFMA, shift/swizzle, transcendental ROM); conversion and load/store type; source, third-source and destination reads and write; mask register reads; operand swaps; scalar or packed; data type; integer-register transfer; rounding; FP flag update.

2.2.2.3. Configuration Parameters

Table 12. Front End configuration parameters

Define	Value	Description
<code>CORE_NR_THREADS</code>	2	Threads per core.
<code>FE_FETCH_BUFFERS</code>	2	Thread-buffer entries per thread.
<code>FE_FETCH_READ_SIZE</code>	256	Bits per fetch block and per entry.
<code>FE_FETCH_PTR_SIZE</code>	5	Read-pointer width: 1 entry bit and 4 halfword bits.
<code>PC_SIZE_EXT</code>	49	Front End PC width.

2.2.3. Pipeline Stages

The Front End pipeline has eight stages, F0 to F7 ([Figure 9](#)). Stages F1 to F5 overlap the L0 micro-cache pipeline.




Figure 9. Front End pipeline

F0 Stage

The F0 stage receives instruction-block requests from both thread buffers and decides which one gets access to the L0 micro-cache using a 2:1 LRU arbiter. The selected request is registered and sent to the L0 micro-cache through the Neighborhood interface. This stage also handles the enabling and disabling of threads and redirect requests from the integer pipeline, for example the new address after a mispredicted branch.

F1–F5 Stage

The selected request is held in F1 until the L0 micro-cache accepts it, which happens at the earliest in the second F1 cycle. The thread then waits for the response. The L0 micro-cache has a fixed five-cycle latency, whether the request is a hit or a miss: counting the cycle in which it accepts the request as the first cycle, the response is returned in the fifth cycle, in F5 (Figure 11).

F6 Stage

The response from the L0 micro-cache is checked. If it is a hit, the instruction block is stored in the buffer of the thread that issued the request. If it is a miss, nothing is stored, and the thread waits for the cache to be filled. Instruction reading is independent of the fetch request: whenever a buffer entry holds instructions, the thread reads its next instruction, expands it if it is compressed, and passes it on for decoding, in parallel with any request in F0 to F5.

F7 Stage

This stage selects which thread's instruction is read from its buffer. The selected instruction is decoded and stored, along with its control signals, in that thread's FIFO. The thread scheduler decides which thread's FIFO issues an instruction to the integer pipeline in the next cycle. If both threads are eligible, they are selected in round-robin order; otherwise, the eligible thread is selected.

With an L0 hit and no competing requests, the first instruction of a block reaches ID nine cycles after its request is issued in F0.

2.2.4. Interaction with the L0 Micro-Cache

Each L0 micro-cache serves the eight threads of four ET-Minion cores. Figure 10 shows the request and response path of one core. This section describes the interface from the Front End side; the Neighborhood side is described in Section 3.2.




Figure 10. Front End to L0 micro-cache request and response path

2.2.4.1. Request Arbitration and Acceptance

A fetch request passes through the following steps:

1. **2:1 LRU arbiter, F0.** Selects between the two threads of the core. With two clients the grant alternates when both threads request. The priority is updated on every grant.
2. **F1 request register.** Holds the granted request on `f1_icache_req` until the Neighborhood accepts it with `f1_icache_req_ready`. The register accepts a new request when it is empty or its current request is accepted in the same cycle; otherwise the 2:1 arbiter issues no grant.
3. **8:1 LRU arbiter, Neighborhood.** The request is first captured in a per-thread request register, then arbitrated against the other seven threads served by the same L0 micro-cache. The grant of this arbiter is returned as `f1_icache_req_ready`. Because of the capture register, a request is accepted at the earliest in its second F1 cycle.

2.2.4.2. Response Timing and Ownership

The L0 micro-cache has a five-stage pipeline and returns exactly one response per accepted request, from its last stage, for a hit and for a miss alike. The response reaches the Front End in F5, four cycles after acceptance; with the earliest acceptance this is five cycles after the request enters F1.

The response carries no thread identifier. Each thread buffer identifies its responses by position in time: when its request is granted, the thread inserts a token (a valid bit and the reserved entry number) into its own F1–F6 tracker. The token remains in F1 until the request is accepted and then advances one stage per cycle. A response registered in F6 together with a valid token belongs to that thread, and the token selects the entry to be written. The response and fill-done signals are registered in `frontend_top` and broadcast to both thread buffers.




Figure 11. Request and hit response

In Figure 11 the request is issued in cycle 1, accepted in cycle 3, and the block returns in cycle 7. It is written in cycle 8; the first instruction is in F7 in cycle 9 and in ID in cycle 10.

2.2.4.3. Miss Handling and Fill Done

A miss response has the same timing as a hit and carries no data. On a miss the thread writes nothing, keeps its fetch PC and enters a sleep state (`f0_miss_pending`) in which it issues no requests.

On completion of a line fill from the L1 instruction cache, or of a translation fill, the L0 micro-cache pulses `f6_icache_fill_done`. The signal is broadcast to all threads of the L0 micro-cache and carries no address. Every sleeping thread wakes and retries; a thread whose line is still absent misses again and sleeps until the next fill. A redirect, a thread disable or a pending halt also clears the sleep state.




Figure 12. Miss, sleep and retry

In Figure 12 the miss response arrives in cycle 7 and the thread sleeps from cycle 9. Fill done arrives in cycle 11, is registered, and the thread retries in cycle 12.

Fill done is issued one cycle after the response slot. A request that looks up the L0 micro-cache while a fill is being written still misses, and its response is the last miss that the fill can cause. Issuing fill done one cycle after that response, and registering it again in the Front End, ensures that every thread has entered the sleep state before the wake-up arrives. A wake-up that arrived first would be lost, and the thread would wait for an unrelated fill.

2.2.4.4. Request Cancellation

A redirect invalidates the thread's tokens in F2 to F5 and suppresses the F6 write, so the responses of in-flight requests are discarded.

A request still waiting in the F1 request register cannot be withdrawn: the register is shared and the Neighborhood already holds a copy. The thread records the cancellation and does not advance the token when the request is accepted. The lookup is performed and its response discarded. The thread's new request is issued only after the cancelled request has been accepted.




Figure 13. Redirect while a request waits in F1

In Figure 13 the redirect arrives in cycle 3 while the request waits in F1. The cancellation is held until the request is accepted in cycle 5; no token enters F2. The new request is issued in cycle 6 and accepted in cycle 8.

2.2.5. Thread Buffer

Each thread has a thread buffer (`frontend_thread_buffer`). It holds the thread's fetch PC and a 64-byte instruction buffer, generates and tracks fetch requests, writes the returned blocks, and extracts one instruction per cycle.




Figure 14. Thread buffer operation




Figure 15. Thread buffer pipeline

2.2.5.1. Buffer Entries

The instruction buffer has two entries of 256 bits (32 bytes). An entry holds one fetch block: eight 32-bit instructions, sixteen compressed instructions, or a combination. Two entries allow the next block to be fetched while the current block is consumed.

The entries are implemented as one latch-based register file with two ports:

- **Write port, 256 bits.** Writes a complete block into one entry, with one write enable per 32-bit word.
- **Read port, 32 bits.** Addressed by the 5-bit read pointer `buffer_ptr`: bit 4 selects the entry, bits 3:0 select a halfword within it. The read returns the 32 bits starting at that halfword.

The read port wraps around: halfword 15 of one entry is followed by halfword 0 of the other. The entries therefore form a ring of 32 halfwords, and a 32-bit instruction that starts in the last halfword of an entry is read in one access.




Figure 16. Thread-buffer entries and read pointer

Table 13. Thread-buffer entry state

State	Description
Empty flag	Entry holds no instruction still to be sent. Set for both entries at reset.
Base address	Address of the request that filled the entry. Defines the words written, the first halfword to extract and the upper PC bits of the entry's instructions.
Fault flags	Page fault and access fault of the block.
Error flags	Bus error and ECC error of the block.
Cacheable flag	Block is in a cacheable region. Used for replay.

2.2.5.2. Program Counter

Each thread buffer holds a 49-bit fetch PC, the address of the next block to request. It is updated as listed in Table 14.

Table 14. Fetch PC update

Event	New fetch PC
Reset, or thread enabled	<code>f0_reset_vector</code> , sign-extended.
Block written into an entry	Start of the next aligned 32-byte block.
Redirect (highest priority)	<code>f0_core_req.pc</code> .

Only the first block after a reset or redirect can start within a 32-byte block; all later blocks are aligned. A miss does not change the fetch PC, so the same address is requested again after the fill.

The PC of each extracted instruction is formed from the upper bits of its entry's base address and the halfword position of the read pointer.

2.2.5.3. Request Generation and Entry Fill

A thread issues a request in F0 when all conditions in Table 15 are met.

Table 15. Request conditions

Condition	Reason
Thread enabled, and not in the enable cycle	Reset vector is loaded at the end of the enable cycle.
No redirect in the current cycle	Redirect PC is loaded at the end of the cycle.
Thread not sleeping	A missed line is requested again only after a fill.
No request of the thread in F1 to F6	At most one request per thread is in flight, which keeps responses in order and allows ownership by timing.
An entry is empty or is released in the current cycle	The response cannot be stalled and must have a destination.
Not in reset; not executing the debug program buffer	—

On a grant, the thread reserves an entry and records the request address as its base address. If both entries are empty, entry 1 is reserved.

Table 16 shows the entry sequence after a thread is enabled at address A. $[m, n]$ gives the block number in entry 0 and entry 1; x denotes an empty entry.

Table 16. Entry fill sequence

Step	Event	Entries
1	Thread enabled at address A.	$[x, x]$
2	Block containing A written into entry 1; extraction starts at the halfword of A.	$[x, A \gg 5]$
3	Entry 0 empty: the next block is requested when the first request completes.	$[(A \gg 5) + 1, A \gg 5]$
4	Last instruction of entry 1 extracted; entry 1 released.	$[(A \gg 5) + 1, x]$
5	Next block requested into entry 1; extraction continues in entry 0.	$[(A \gg 5) + 1, (A \gg 5) + 2]$

A block write enables only the words from the word containing the request address to word 7. The register file requires its write-data enables one clock phase before the write; the thread computes them in F5 from its token and `f5_icache_resp_miss`, holds them in a phase latch, and applies the word write enables in F6.

2.2.5.4. Instruction Extraction

When both entries are empty, the read pointer is loaded in F5 with the halfword position of the request address within the reserved entry, so that extraction can start within a block. Afterwards the pointer advances with each extracted instruction.

An instruction is available when the entry it is read from is not empty. A 32-bit instruction that starts in halfword 15 requires both entries to be non-empty; the wrap-around read returns both halves, the pointer moves to halfword 1 of the other entry, and the first entry is released.

If the entry being read is empty but its block is written in the same cycle, the instruction is taken directly from the response data (bypass). The upper half of an instruction that crosses into an entry being written is bypassed in the same way. The first instruction of a block is therefore extracted in the cycle the block is written. The bypass does not check the destination entry of the response; with one request in flight per thread, the only block that can arrive is the one being waited for.

2.2.5.5. RVC Expander

The 32 bits read at the pointer are passed to the RVC expander (`frontend_rvc_expander`), a combinational RV64C decompressor:

- If bits [1:0] are not **11**, the instruction is compressed. It is expanded to its 32-bit equivalent and the pointer advances by 16 bits.
- If bits [1:0] are **11**, the instruction is 32 bits wide. It passes through unchanged and the pointer advances by 32 bits.

All instructions are therefore 32 bits wide after F6. The `rvc` bit of `frontend_core_resp` indicates a compressed instruction; the integer pipeline uses it to advance the PC by 2 instead of 4.

2.2.5.6. Fault and Error Status

Each instruction carries the status of the entries it was read from:

Table 17. Instruction fetch status

Field	Source
<code>pf0, af0</code>	Fault flags of the entry holding the start of the instruction.
<code>pf1, af1</code>	Fault flags of the entry holding the upper half of an instruction that crosses entries; zero otherwise.
<code>bus_error, ecc_error</code>	Error flags of the instruction's entry, or of both entries for an instruction that crosses entries.
<code>replay</code>	Set when the entry's block is not cacheable and <code>f0_core_req.speculative</code> indicates older instructions of the thread in flight. Captured when the instruction is written into the FIFO.

All status fields are forced to zero while the debug program buffer is executed.

2.2.6. Entry Invalidation, Redirect and Thread Start

2.2.6.1. Entry Invalidation

An entry is released when its last instruction leaves F6, or when the whole buffer is flushed ([Table 18](#)).

Table 18. Entry release conditions

Condition	Entries released
32-bit instruction at byte offset 28 of the entry is extracted	This entry

Condition	Entries released
Compressed instruction at byte offset 30 is extracted	This entry
32-bit instruction at byte offset 30 is extracted (upper half in the other entry)	This entry
Redirect (<code>f0_core_req_valid</code>) or thread disable	Both entries
Entry to debug mode	Both entries

A response that reports a fault or error marks the reserved entry non-empty even if no data is written, so that the fault is delivered to the integer pipeline with the next instruction read from that entry.

2.2.6.2. Redirect

The integer pipeline redirects a thread by asserting `f0_core_req_valid` with the new PC in `f0_core_req.pc`. Redirects are issued for branch and jump mispredictions, exceptions and returns from trap, replays, inter-processor interrupts, and entry to or exit from debug mode. Table 19 lists the effect in the redirect cycle.

Table 19. Effect of a redirect

Stage	Effect
F0	Fetch PC loaded with the redirect PC. No request in this cycle. Sleep state cleared.
F1	Request waiting in the F1 request register marked cancelled.
F2–F5	Tokens invalidated.
F6	No block written. Both entries released.
F7	Instruction discarded.
FIFO	Emptied. The thread is not eligible for scheduling in this cycle.

The target block is requested in the following cycle. Its first instruction reaches ID ten cycles after the redirect (Figure 17).




Figure 17. Redirect to first target instruction

2.2.6.3. Thread Enable and Reset Vector

`f0_thread_enabled` is registered on entry to the Front End, and one-cycle enable and disable events

are derived from its edges.

- **Enable.** The fetch PC is loaded with `f0_reset_vector` at the end of the enable cycle, and the first request is issued in the next cycle. The first instruction reaches ID eleven cycles after `f0_thread_enabled` rises (Figure 18).
- **Disable.** No further requests are issued, both entries are released and the sleep state is cleared. The instruction in F7 and the FIFO contents are not cleared; the thread scheduler no longer selects the thread.




Figure 18. Thread enable to first instruction

After reset, both entries are empty and the fetch PC holds the reset vector.

2.2.7. Decoders

One integer decoder (`intpipe_decode`) and one VPU decoder (`vpu_decoder`) are shared by both threads (Figure 19). In F7, the instruction of one thread passes through both decoders in parallel, and the instruction and both control words are written into that thread's FIFO. Decoding before the FIFO removes decode logic from the ID stage.




Figure 19. Decoders, instruction FIFOs and thread scheduler

Both decoders are combinational look-up tables over the 32-bit instruction:

- **Integer decoder.** Produces `minion_control` (Table 11). An unlisted encoding is marked not legal, and the integer pipeline raises an illegal-instruction exception.
- **VPU decoder.** Produces `vpu_ctrl_sigs_t` and a flag that identifies VPU instructions. An unlisted encoding is a non-VPU instruction.

A thread is a decode candidate when it has an extracted instruction, its FIFO is not full and it is awake (Section 2.2.9). The decoded thread is selected as follows:

Table 20. Decode thread selection

Candidates	Thread decoded
Thread 0 only	Thread 0
Thread 1 only	Thread 1
Both	Thread not presented in ID in the current cycle
None	Thread not presented in ID in the current cycle

The instruction is written into the FIFO when the selected thread has an extracted instruction and its FIFO is not full; the awake condition affects only the selection.

2.2.8. Instruction FIFO

Each thread has a two-entry instruction FIFO between F7 and ID. It decouples fetch from issue: while ID is stalled or serving the other thread, the thread continues to extract and decode, and the decoded instructions are held in the FIFO. An entry is a `frontend_thread_data` (Table 10): the instruction, its fetch status and both control words.

The FIFO is written when its thread is selected for decode and read when the integer pipeline accepts the thread's instruction (`id_inst_valid`, `id_inst_ready` and `id_inst_thread_id`). It uses 2-bit read and write pointers; it is full when the pointers differ only in the wrap bit and empty when they are equal.

The storage is organized for timing rather than as an indexed array. Entry 0 is always the head, whose content drives the ID outputs; entry 1 is a holding register. The ID multiplexer therefore selects between two fixed registers, one per thread.

Table 21. Instruction FIFO update

Occupancy	Event	Update
Empty	Write	Entry 0 ← new instruction
One	Write and read	Entry 0 ← new instruction
One	Write	Entry 1 ← new instruction
One	Read	None; FIFO becomes empty
Full	Read	Entry 0 ← entry 1

The VPU control word of an entry is written only for VPU instructions. For other instructions it retains the value of an earlier VPU instruction, so `id_vpu_decoder_sigs` and `id_vpu_core_ctrl` are valid only for VPU instructions.

A redirect empties the thread's FIFO.

2.2.9. Thread Scheduler

The integer pipeline accepts one instruction per cycle from either thread. The thread scheduler (`frontend_thread_sched`) selects the thread whose FIFO head is presented to ID in the next cycle.

Awake

The thread is enabled and `id_core_stall` is deasserted. Awake is registered, so a stall affects scheduling one cycle after it is asserted.

Ready

The thread's FIFO is not empty or is being written in the current cycle, and the thread is not being redirected. Including the write allows an instruction to pass from F7 to ID in consecutive cycles.

A thread that is awake and ready is eligible. The scheduler selects between eligible threads in round-robin order:

Table 22. Thread selection

Thread 0 eligible	Thread 1 eligible	Selected
Yes	No	Thread 0
No	Yes	Thread 1
Yes	Yes	Thread not selected previously
No	No	Thread 0

The previous selection is updated only when at least one thread is eligible. The selection is registered and drives `id_inst_thread_id` in the next cycle.

The head of the selected thread's FIFO drives `id_inst_data`, `id_intpipe_ctrl`, `id_vpu_decoder_sigs` and `id_vpu_core_ctrl`. `id_inst_valid` is asserted when that FIFO is not empty, and the integer pipeline accepts the instruction with `id_inst_ready`. `id_inst_valid` is not qualified with `id_core_stall`; the integer pipeline applies the stall. The thread in ID is reselected every cycle, so an instruction that is not accepted can be replaced by the other thread's instruction in the next cycle; it remains at the head of its FIFO.




Figure 20. Round-robin issue with a stall

In Figure 20 both threads are eligible and alternate. `id_core_stall` of thread 1 is asserted in cycles 3 and 4, and thread 1 is not awake in cycles 4 and 5. Thread 0 is selected in cycles 3 to 5 and

presented in ID in cycles 4 to 6; the selection alternates again from cycle 6.

2.2.10. Halt and Debug Program Buffer

2.2.10.1. Halt

The integer pipeline takes a halt on a valid instruction in ID. While `halt` is asserted for a thread, the thread buffer generates a placeholder instruction instead of fetching:

1. The thread follows the normal request conditions but does not bid for the request port. The request is treated as accepted internally, and its token enters the tracker.
2. The token passes through F1 to F6.
3. In F6 an entry write is forced without a response. The data is the last registered response data; fault and error flags are cleared, and the instruction is treated as 32 bits wide.
4. The placeholder passes through decode and the FIFO to ID, where the integer pipeline takes the halt.

The content of the placeholder is undefined. A pending halt also clears the sleep state.

2.2.10.2. Debug Program Buffer

In debug mode, entry 0 of the halted thread's buffer is used as the program buffer: the debugger writes instructions into it and the thread executes them.

Table 23. Debug program-buffer sequence

Step	Operation
Entry	A redirect with <code>debug_info.halt</code> arrives; the thread stops fetching in the same cycle. In the next cycle both entries are released, the extracted instruction is discarded, and words 0 to 4 of entry 0 are written with <code>EBREAK</code> .
Write	The debugger writes the program buffer through <code>debug_ffb_wdata</code> , <code>debug_ffb_en</code> and <code>debug_ffb_thread_sel</code> (Table 24). The words are written one cycle after the enables.
Execute	<code>debug_ffb_exec</code> resets the read pointer to halfword 0 of entry 0. In the next cycle entry 0 is marked non-empty and entry 1 empty, and instructions are extracted from word 0.
Exit	A redirect with <code>debug_info.resume</code> ends program-buffer mode; fetch restarts at the redirect PC.

Table 24. Program-buffer word mapping

Word	Write enable	Data
0	<code>debug_ffb_en[0]</code>	<code>debug_ffb_wdata[31:0]</code>
1	<code>debug_ffb_en[1]</code>	<code>debug_ffb_wdata[63:32]</code>
2	<code>debug_ffb_en[2]</code>	<code>debug_ffb_wdata[31:0]</code>
3	<code>debug_ffb_en[3]</code>	<code>debug_ffb_wdata[31:0]</code>

Word	Write enable	Data
4	—	EBREAK

Word 4 cannot be written by the debugger and always holds **EBREAK**, so execution returns to debug mode at the end of the program buffer. While the program buffer is executed, all fault, error and replay status is forced to zero. A debug reset while the thread is halted rewrites words 0 to 4 with **EBREAK**.

2.3. Integer Pipeline

2.3.1. Overview

Figure: Integer pipeline block diagram.

2.3.2. Pipeline Stages

Figure: Integer pipeline stages (ID, EX, TAG, MEM, WB).

2.3.3. Interfaces

Table: Integer pipeline interface signals.

2.3.4. Register File and Operand Read

2.3.5. Immediate Generation and ALU

2.3.6. Multiply/Divide Unit

2.3.7. Scoreboards and Hazards

2.3.7.1. Integer Scoreboard

2.3.7.2. FP Scoreboard

2.3.7.3. Mask Scoreboard

2.3.8. Hart Interleaving and Stalls

2.3.9. Branches, Redirects and Replay

2.3.10. Interface to the Data Cache

2.3.11. Interface to the VPU

2.3.12. CSRs and Privilege Modes

Table: CSR groups.

2.3.13. Exceptions and Interrupts

2.3.14. Debug and Breakpoints

2.3.15. Performance Counters

2.3.16. Verification Notes

2.4. Data Cache

2.4.1. Overview

Figure: Data cache block diagram.

2.4.2. Organization and Operating Modes

Table: Data cache operating modes.

2.4.3. Interfaces

Table: Data cache interface signals.

2.4.4. Request Pipeline

Figure: Data cache request pipeline.

2.4.5. Address Translation and PMA Checks

2.4.6. Tag, Data and LRU Arrays

2.4.7. Miss Handling

2.4.8. Replay Queue

2.4.9. Writeback and Store Merge

2.4.10. Atomic Operations

2.4.11. Cache-Management Operations

2.4.12. Scratchpad Control

2.4.13. Tensor Load and Reduce

2.4.14. Interface to the Neighborhood

2.4.15. Errors

2.4.16. Performance

2.4.17. Verification Notes

2.5. Vector Processing Unit

2.5.1. Overview

Figure: VPU block diagram.

2.5.2. Interfaces

Table: VPU interface signals.

2.5.3. Control and Micro-Instruction Decode

2.5.4. Register Files

2.5.5. Lanes

Figure: VPU lane.

2.5.6. Mask Unit

2.5.7. Bypass Network

2.5.8. Shift and Swizzle Unit

2.5.9. TxFMA Unit

2.5.10. Integer Multiply-Accumulate Unit

2.5.11. Transcendental Unit

2.5.12. Tensor Operations

2.5.12.1. Tensor FMA

2.5.12.2. Tensor Quantization

2.5.12.3. Tensor Reduce

2.5.13. Interfaces to the Integer Pipeline and Data Cache

2.5.14. Performance

2.5.15. Verification Notes

3. Neighborhood

3.1. Neighborhood Overview

3.1.1. Partitioning

Figure: Neighborhood block diagram.

3.1.2. Shared Resources and Core Connections

3.1.3. Voltage-Domain Crossings

3.1.4. External Interfaces

Table: Neighborhood interface groups.

3.2. Front End – I-Cache Interface

3.2.1. Overview

Figure: Front End to I-cache connection within a Neighborhood.

3.2.2. Interface Signals

Table: Front End – I-cache interface signals.

3.2.3. Request Arbitration

3.2.4. Request and Response Timing

Figure: Timing of a request and its hit response.

3.2.5. Miss, Sleep and Fill-Done Protocol

Figure: Timing of a miss, fill and retry.

3.2.6. Fault and Error Reporting

3.2.7. Livelock Avoidance

3.2.8. Known Limitations

3.3. L0 Micro-Cache

3.3.1. Overview

Table: L0 micro-cache characteristics.

3.3.2. Internal Blocks

Figure: L0 micro-cache block diagram.

3.3.3. Pipeline Stages

Figure: L0 micro-cache pipeline.

3.3.4. Instruction TLB

3.3.5. Micro-Tag and Micro-Data Arrays

3.3.6. PMA Check

3.3.7. Replacement

3.3.8. Miss Handling

3.3.8.1. TLB Miss

3.3.8.2. Tag Miss

3.3.8.3. Fill and Thread Wake-Up

3.3.9. L1 Bypass

3.3.10. Software Prefetch

Figure: Prefetch state machine.

3.4. L1 Instruction Cache

3.4.1. Overview

Figure: L1 instruction cache block diagram.

3.4.2. Organization

3.4.3. Tag, LRU and Data Arrays

3.4.4. External Data RAM

3.4.5. Request Arbitration Between L0 Micro-Caches

3.4.6. Miss Path

3.4.7. Errors and ECC

3.5. Shared Page-Table Walker

3.5.1. Clients and Arbitration

3.5.2. Walk Sequence and Page Sizes

Figure: Page-table walk state machine.

3.5.3. PTE Cache and Invalidation

3.6. Memory Request and Response Path

3.6.1. Miss, Evict and Fill FIFOs

Figure: Neighborhood memory request and response path.

3.6.2. Request Arbitration

3.6.3. Response Routing and Backpressure

3.7. Tensor Support

3.7.1. Cooperative Tensor Load

3.7.2. Tensor Store Buffer

3.7.3. TBOX Router

3.8. Synchronization, Messaging, Registers and PMU

3.8.1. Fast Local Barrier

3.8.2. Local Message Network

3.8.3. IPI and FCC Paths

3.8.4. Neighborhood Registers and APB Access

3.8.5. Debug Channel

3.8.6. Performance Monitoring Unit